



Domain-Specific AI Assistant for the UL FRI Student Office

Aleksa Sibinović, Hristijan Milanovski, and Sara Ivanovska

Abstract

The abstract goes here.

Keywords

Retrieval-Augmented Generation, Domain-Specific Assistant, Large Language Models, Question Answering, Student Office

Advisors: Slavko Žitnik

Introduction

Dataset

Since no pre-existing dataset exists for this domain, we construct our own knowledge base from publicly available institutional sources.

Data Sources

We collect data from three main sources: the **UL FRI website** (fri.uni-lj.si), covering study program descriptions, enrollment procedures, exam rules, and the academic calendar; the **UL central website** (uni-lj.si), containing university-wide student regulations and administrative guidelines; and **official PDF documents**, such as the Student Statute and enrollment rulebooks.

Data Collection and Processing

Web pages will be scraped and cleaned to retain only the main textual content, discarding navigational elements and footers. Each document will be stored alongside metadata including its source URL and retrieval date. PDF documents will be processed using text extraction libraries to produce clean plain-text representations. The collected documents will then be split into overlapping chunks to enable fine-grained retrieval of relevant passages, following standard RAG preprocessing practices.

References

Methods

Results

Discussion